

Machine Learning Project

Binary classification based on 3 layers neural network (2)

First layer

$Z^{[1]} = W^{[1]}X + b^{[1]}$: X denotes the input data

$A^{[1]} = g^{[1]}(Z^{[1]})$: $g^{[1]}$ is the activation function at the first layer

Second layer

$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]}$

$A^{[2]} = g^{[2]}(Z^{[2]})$: $g^{[2]}$ is the activation function at the second layer

Third layer

$Z^{[3]} = W^{[3]}A^{[2]} + b^{[3]}$

$A^{[3]} = g^{[3]}(Z^{[3]})$: $g^{[3]}$ is the activation function at the third (output) layer

Activation Function

- Sigmoid

$$g(z) = \frac{1}{1 + \exp^{-z}}$$

- tanh

$$g(z) = \frac{\exp^z - \exp^{-z}}{\exp^z + \exp^{-z}}$$

- ReLU

$$g(z) = \max(0, z)$$

- Leaky ReLU

$$g(z) = \max(\alpha z, z), \quad \alpha \in \mathbb{R}^+$$

Neural Network Architecture

- The sizes of the hidden layers and the output layer should be determined with respect to the validation accuracy obtained by the network architecture with all the activation functions being sigmoid functions. (

$$g^{[1]} = g^{[2]} = g^{[3]} = \text{Sigmoid}$$

- Apply different activation functions at all the layers except the output layer that should be Sigmoid function
- Apply different activation functions at different layers except the output layer that should be Sigmoid function

Dataset

- The dataset consists of human images and horse images for the training and the validation
- The classifier should be trained using the training set
- The classifier should be tested using the validation set
- Vectorize an input image matrix into a column vector

Implementation

- Write codes in python programming
- Use jupyter notebook for the programming environment
- You can use any library
- You have to write your own functions for the followings:
 - compute the forward propagation
 - compute the backward propagation
 - compute the loss
 - compute the accuracy
 - compute the gradient of the model parameters with respect to the loss
 - update the model parameters
 - plot the results

Optimization

- Apply the gradient descent algorithm with an appropriate learning rate
- Apply the number of iterations that lead to the convergence of the algorithm
- Use the vectorization scheme in the computation of gradients and the update of the model parameters

[Implementation]

(1) Libraries and Global variables

```
In [3]:  
  
from torch.utils.data import Dataset, DataLoader  
import torchvision.transforms as transforms  
import matplotlib.pyplot as plt  
import numpy as np  
import torchvision  
import torch  
import math  
import os  
  
# Global Variables  
train_data_path = './horse-or-human/train'  
validation_data_path = './horse-or-human/validation'  
  
layer_dims = [10000,50,10,1]      # number of units(Neurons) in each layer  
learning_rate = 0.02               # step size per each epoch (iteration)  
threshold = 0.1                   # minimum of cost  
max_epoch = 2500                  # maximum number of epoch (iteration)
```

(2) Generate Input matrix X and Output vector Y from training and validation datasets

```
In [4]:  
  
def initialize_inputs(image_path) :  
    transform = transforms.Compose([transforms.Grayscale(),transforms.ToTensor(),])  
    # the code transforms.Grayscale() is for changing the size [3,100,100] to [1, 100, 100]  
    # (notice : [channel, height, width] )  
    image_set = torchvision.datasets.ImageFolder(root=image_path, transform=transform)  
    loader = torch.utils.data.DataLoader(image_set, batch_size=1, shuffle=False, num_workers=1)  
  
    for i,data in enumerate(loader) :  
        image, label = data  
        image = image.view(10000,1)  
        label = label.view(1,1).type(torch.FloatTensor)  
  
        if i == 0 :  
            t_images = image  
            t_labels = label  
        else :  
            t_images = torch.cat((t_images,image),dim = 1)  
            t_labels = torch.cat((t_labels,label),dim = 1)  
  
        images = t_images.numpy()  
        labels = t_labels.numpy()  
  
    return images, labels
```

(3) Activation Functions

```
In [31]:  
  
def Sigmoid(Z):  
    A = 1/(1+np.exp(-Z))  
    return A  
  
def Sigmoid_backward(dA, Z):  
    t_A = 1/(1+np.exp(-Z))  
    dZ = dA * t_A * (1-t_A)  
    return dZ  
  
def ReLU(Z):  
    A = np.maximum(0,Z)  
    return A  
  
def ReLU_backward(dA, Z):  
    dZ = np.array(dA, copy=True)  
    dZ[Z <= 0] = 0  
    return dZ  
  
def Leaky_ReLU(Z) :  
    A = np.maximum(0.01*Z, Z)  
    return A  
  
def Leaky_ReLU_backward(dA,Z) :  
    dZ = np.array(dA, copy=True)  
    dZ[Z < 0] *= 0.01  
    return dZ  
  
def Tanh(Z) :  
    A = np.tanh(Z)  
    return A  
  
def Tanh_backward(dA, Z) :  
    t_A = np.tanh(Z)  
    dZ = dA * (1 - np.square(t_A))  
    return dZ
```

(4) Cost and Parameter initialization / update Functions

```

In [20]:
def initialize_parameters(n) :
    parameters = dict()

    # Initialize W[i] and b[i] for i in [1,L-1]
    for l in range(1, len(n)) :
        parameters['W'+str(l)] = np.random.randn(n[l],n[l-1]) / np.sqrt(n[l-1])
        parameters['b'+str(l)] = np.zeros((n[l],1))

    return parameters

def update_parameters(parameters,gradients,learning_rate) :
    L = len(parameters) // 2

    # Update W[i] and b[i] for i in [1,L]
    for l in range(1,L+1) :
        dW, db = gradients['dW'+str(l)], gradients['db'+str(l)]
        parameters['W'+str(l)] -= learning_rate * dW
        parameters['b'+str(l)] -= learning_rate * db

    return parameters

def cost_computation(AL, Y) :
    m = Y.shape[1]

    cost = (-np.dot(Y,np.log(AL).T) - np.dot(1-Y,np.log(1-AL).T)) / m
    cost = np.squeeze(cost)

    return cost

```

(5) Forward Propagation

```

In [21]:
def forward_Z_computation(A_prev,W,b) :
    Z = np.dot(W,A_prev) + b
    return Z

def forward_A_computation(A_prev,W,b,activation) :
    assert activation in ['sigmoid','relu','tanh','leaky_relu']

    Z = forward_Z_computation(A_prev,W,b)
    if activation == 'sigmoid' :
        A = Sigmoid(Z)
    elif activation == 'relu' :
        A = ReLU(Z)
    elif activation == 'tanh' :
        A = Tanh(Z)
    else :
        A = Leaky_ReLU(Z)

    cache = ((A_prev,W,b),Z)
    return A, cache

def forward_propagation(X,parameters,activations) :
    caches = []
    A = X
    L = len(parameters) // 2

    for l in range(1,L+1) :
        A_prev, W, b = A, parameters['W'+str(l)], parameters['b'+str(l)]
        A,cache = forward_A_computation(A_prev,W,b,activations[l-1])
        caches.append(cache)
    AL = A

    return AL, caches

```

(6) Backward Propagation

```

In [22]:
def backward_params_dev_computation(dZ, cache) :
    A_prev, W, b = cache
    m = A_prev.shape[1]

    dW = np.dot(dZ,a_prev.T) / m
    db = np.sum(dZ, axis = 1, keepdims = True) / m
    dA_prev = np.dot(W.T,dZ)

    return dA_prev, dW, db

def backward_params_dev_computation(dA,cache,activation) :
    assert activation in ['sigmoid','relu','tanh','leaky_relu']
    (A_prev,W,b),Z = cache
    m = A_prev.shape[1]

    if activation == 'sigmoid' :
        dZ = Sigmoid_backward(dA,Z)
    elif activation == 'relu' :
        dZ = ReLU_backward(dA,Z)
    elif activation == 'tanh' :
        dZ = Tanh_backward(dA,Z)
    else :
        dZ = Leaky_ReLU_backward(dA,Z)

    dW = np.dot(dZ,A_prev.T) / m
    db = np.sum(dZ, axis = 1, keepdims = True) / m
    dA_prev = np.dot(W.T,dZ)

    return dA_prev, dW, db

def backward_propagation(AL, Y, caches,activations) :
    gradients = dict()
    L = len(caches)
    m = AL.shape[1]
    Y = Y.reshape(AL.shape)

    dAL = - (np.divide(Y,AL) - np.divide(1-Y,1-AL))
    cache = caches[L-1]
    dA_prev, dW, db = backward_params_dev_computation(dAL,cache,activations[L-1])
    gradients['dA'+str(L)] = dA_prev
    gradients['dW'+str(L)] = dW
    gradients['db'+str(L)] = db

    for l in reversed(range(1,L)) :
        dA = gradients['dA' + str(l+1)]
        dA_prev, dW, db = backward_params_dev_computation(dA,caches[l-1],activations[l-1])
        gradients['dA'+str(l)] = dA_prev
        gradients['dW'+str(l)] = dW
        gradients['db'+str(l)] = db

    return gradients

```

(7) Predictions

In [9]:

```
def predict(AL,Y) :
    correct = np.zeros(Y.shape)
    m = AL.shape[1]

    prediction = AL > 0.5
    correct = prediction == Y

    accuracy = np.sum(correct) / m
    return accuracy
```

(8) 3-layer Neural Network

In [43]:

```
def NN_3_layers(X,Y,t_X,t_Y,n,activations,learning_rate,threshold,max_epoch) :

    np.random.seed(1)
    costs, t_costs = [],[]
    accuracy, t_accuracy = [],[]
    parameters = initialize_parameters(n)

    for epoch in range(max_epoch) :
        AL, caches = forward_propagation(X, parameters, activations)
        t_AL, _t = forward_propagation(t_X, parameters, activations)

        cost = cost_computation(AL, Y)
        t_cost = cost_computation(t_AL, t_Y)

        gradients = backward_propagation(AL, Y, caches,activations)

        parameters = update_parameters(parameters,gradients,learning_rate)

        costs.append(cost)
        t_costs.append(t_cost)

        accuracy.append(predict(AL,Y))
        t_accuracy.append(predict(t_AL,t_Y))

    #     if epoch % 10 == 0:
    #         print ("Cost after iteration %i: %f" %(epoch, cost))
    if cost < threshold :
        break

    return costs, accuracy, t_costs, t_accuracy
```

(9) Plot Loss and Accuracy Graphs

In [11]:

```
def plotResult(costs,accuracy,title) :
    plt.rcParams["figure.figsize"] = (12,6)
    plt.subplot(121)
    plt.title('Costs', fontsize = 15, color = 'black')
    plt.plot(np.squeeze(costs))
    plt.subplot(122)
    plt.title('Accuracy', fontsize = 15, color = 'black')
    plt.plot(np.squeeze(accuracy))
    plt.suptitle(title, fontsize = 20, color = 'black',position=(0.5, 1.0+0.05))
    plt.show()
```


Output

- Plot the training loss at every iteration (x-axis: iteration, y-axis: loss)
- Plot the validation loss at every iteration (x-axis: iteration, y-axis: loss)
- Plot the training accuracy at every iteration (x-axis: iteration, y-axis: accuracy)
- Plot the validation accuracy at every iteration (x-axis: iteration, y-axis: accuracy)
- Present the table for the final accuracy and loss with training and validation datasets with your best neural network architecture as below:

```
In [14]:  
X,Y = initialize_inputs(train_data_path)  
t_X, t_Y = initialize_inputs(validation_data_path)  
  
n = [10000,50,10,1]
```

(1) $g^{[1]}, g^{[2]}, g^{[3]}$ are Sigmoid (from the previous assignment)

- Learning curves
- Loss and Accuracy table


dataset	loss	accuracy
training	0.48654	80.82 %
validation	0.39640	87.50 %

```
In [61]:  
activations = ['sigmoid', 'sigmoid', 'sigmoid']  
learning_rate = 0.01  
thres_hold = 0.01  
max_epoch = 3000  
  
costs, accuracy, t_costs, t_accuracy = NN_3_layers(X,Y,t_X,t_Y,n,activations,learning_rate,threshold,max_epoch)
```

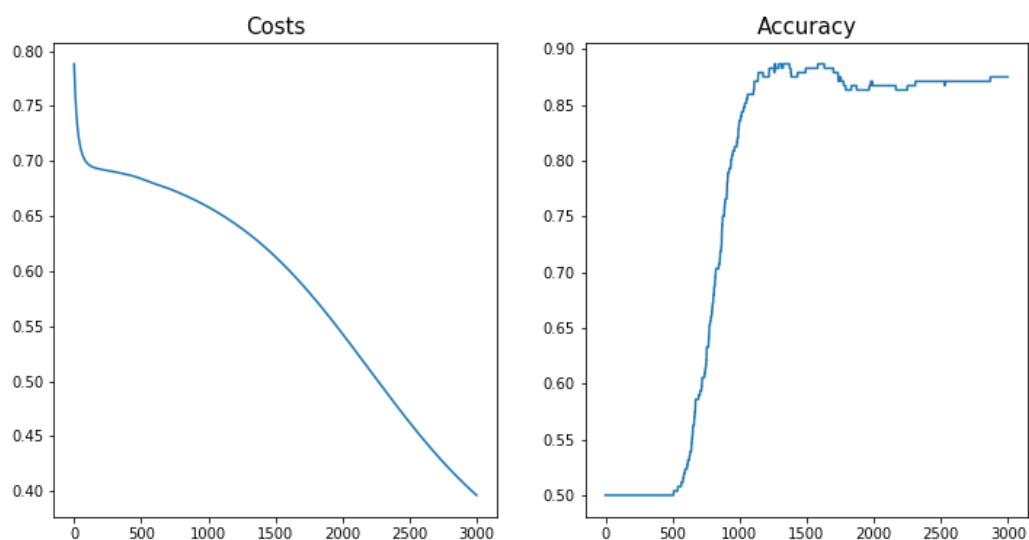
In [62]:

```
plotResult(costs, accuracy, 'Training')  
plotResult(t_costs, t_accuracy, 'Validation')
```

Training



Validation



In [63]:

```
print("Training Loss      : %.5f" % costs[-1])  
print("Training Accuracy  : %.2f %%" % (accuracy[-1] * 100))  
print("Validation Loss     : %.5f" % t_costs[-1])  
print("Validation Accuracy : %.2f %%" % (t_accuracy[-1] * 100))
```

```
Training Loss      : 0.48654  
Training Accuracy  : 80.82 %  
Validation Loss     : 0.39640  
Validation Accuracy : 87.50 %
```

(2) $g^{[1]}, g^{[2]}$ are tanh and $g^{[3]}$ is Sigmoid

- Learning curves
- Loss and Accuracy table

dataset	loss	accuracy
training	0.22375	95.33 %
validation	0.37413	81.25 %

In [44]:

```

activations = ['tanh', 'tanh', 'sigmoid']
learning_rate = 0.002
thres_hold = 0.01
max_epoch = 2500

```

```
costs, accuracy, t_costs, t_accuracy = NN_3_layers(X,Y,t_X,t_Y,n,activations,learning_rate,threshold,max_epoch)
```

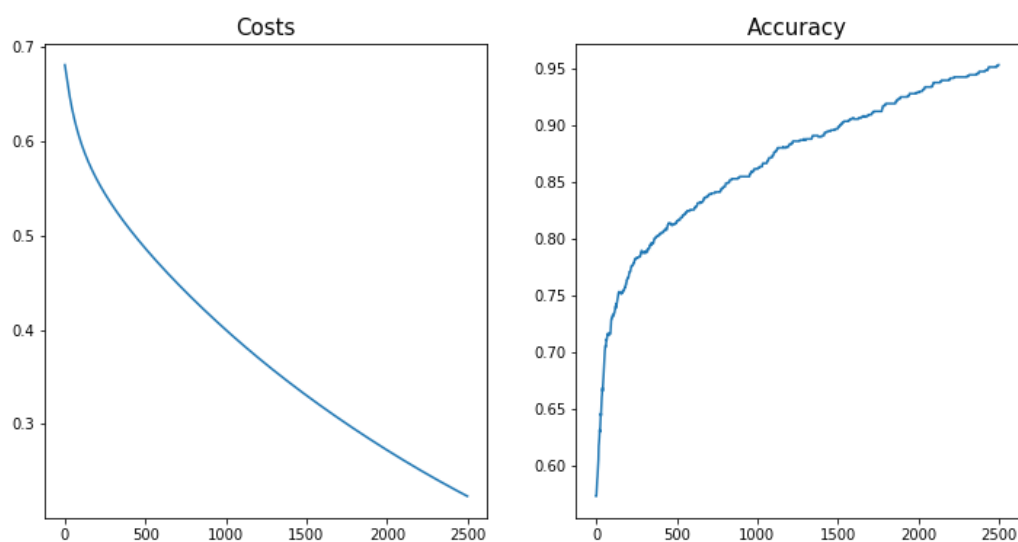
In [45]:

```

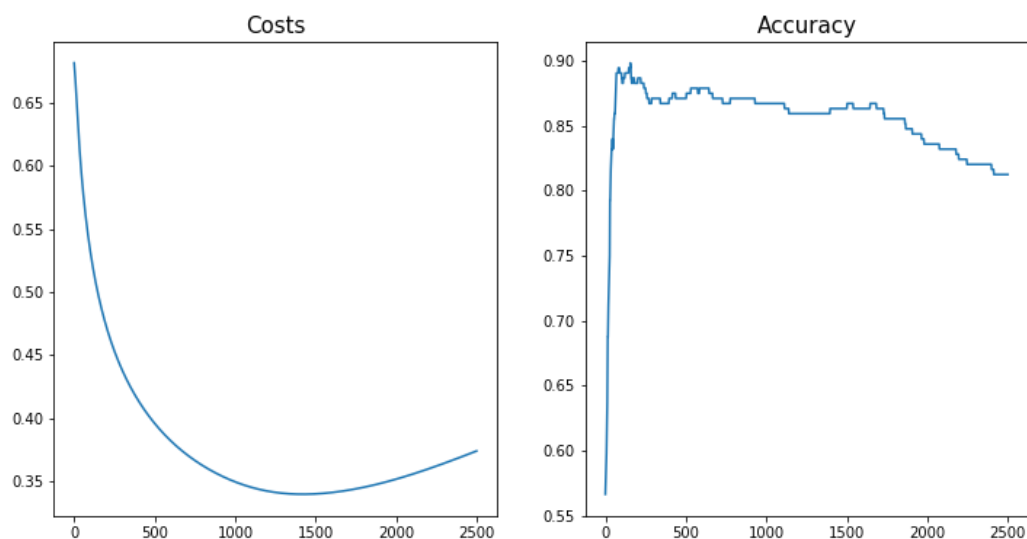
plotResult(costs, accuracy, 'Training')
plotResult(t_costs, t_accuracy, 'Validation')

```

Training



Validation



In [46]:

```

print("Training Loss      : %.5f" % costs[-1])
print("Training Accuracy  : %.2f %% " % (accuracy[-1] * 100))
print("Validation Loss    : %.5f" % t_costs[-1])
print("Validation Accuracy: %.2f %% " % (t_accuracy[-1] * 100))

```

```

Training Loss      : 0.22375
Training Accuracy  : 95.33 %
Validation Loss    : 0.37413
Validation Accuracy: 81.25 %

```

(3) $g^{[1]}, g^{[2]}$ are ReLU and $g^{[3]}$ is Sigmoid

- Learning curves
- Loss and Accuracy table

dataset	loss	accuracy
training	0.18629	95.33%
validation	0.40157	80.86%

In [45]:

```

activations = ['relu', 'relu', 'sigmoid']
learning_rate = 0.002
thres_hold = 0.01
max_epoch = 2500


```

```
costs, accuracy, t_costs, t_accuracy = NN_3_layers(X,Y,t_X,t_Y,n,activations,learning_rate,threshold,max_epoch)
```

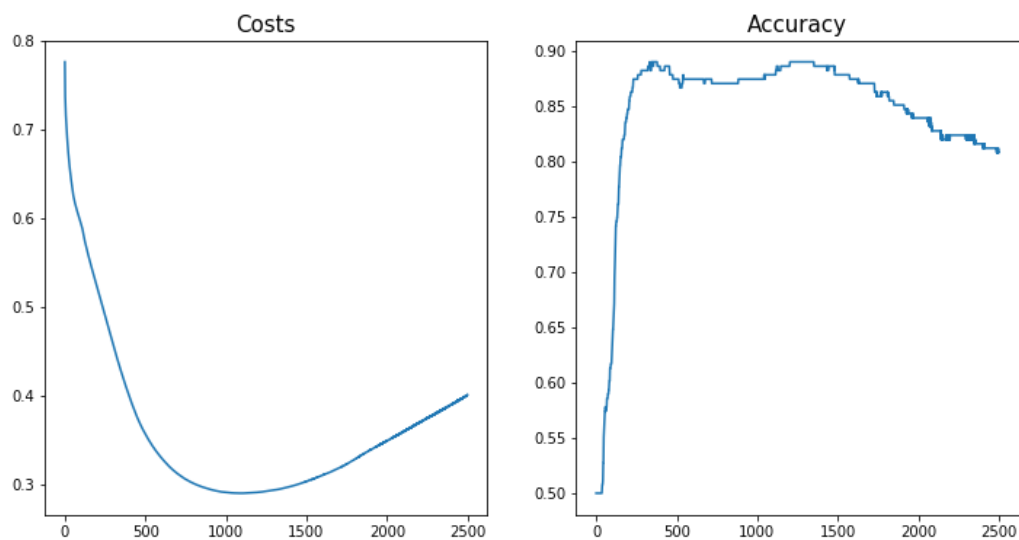
In [46]:

```
plotResult(costs, accuracy, 'Training')
plotResult(t_costs, t_accuracy, 'Validation')
```

Training



Validation



In [52]:

```
print("Training Loss      : %.5f" % costs[-1])
print("Training Accuracy  : %.2f %" % (accuracy[-1] * 100))
print("Validation Loss     : %.5f" % t_costs[-1])
print("Validation Accuracy : %.2f %" % (t_accuracy[-1] * 100))
```

```
Training Loss      : 0.18629
Training Accuracy  : 95.33 %
Validation Loss     : 0.40157
Validation Accuracy : 80.86 %
```

(4) $g^{[1]}, g^{[2]}$ are Leaky ReLU ($\alpha = 0.1$) and $g^{[3]}$ is Sigmoid

- Learning curves
- Loss and Accuracy table

dataset	loss	accuracy
training	0.18501	95.23 %
validation	0.40629	79.69 %

In [35]:

```

activations = ['leaky_relu', 'leaky_relu', 'sigmoid']
learning_rate = 0.002
thres_hold = 0.01
max_epoch = 2500

```

```
costs, accuracy, t_costs, t_accuracy = NN_3_layers(X,Y,t_X,t_Y,n,activations,learning_rate,threshold,max_epoch)
```


In [36]:

```

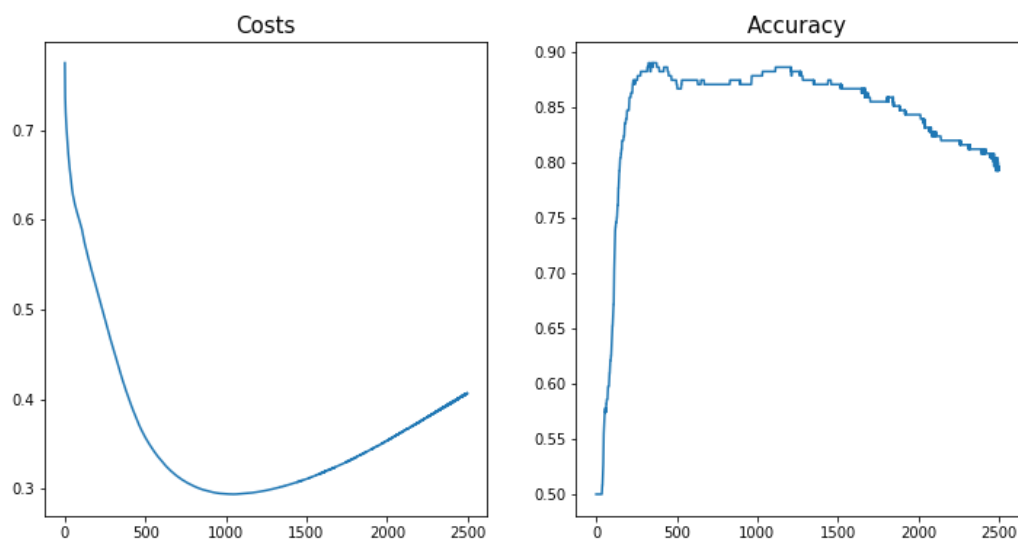
plotResult(costs, accuracy, 'Training')
plotResult(t_costs, t_accuracy, 'Validation')

```

Training



Validation



```
In [37]:
```

```
print("Training Loss      : %.5f" % costs[-1])  
print("Training Accuracy  : %.2f %%" % (accuracy[-1] * 100))  
print("Validation Loss    : %.5f" % t_costs[-1])  
print("Validation Accuracy : %.2f %%" % (t_accuracy[-1] * 100))
```

```
Training Loss      : 0.18501  
Training Accuracy  : 95.23 %  
Validation Loss    : 0.40629  
Validation Accuracy : 79.69 %
```